# How to Choose a Good LLM

Team Members:
- Wang, Yifan
- Khanna, Sunreet
- Li, Jason
- Xu, Jiabao

# Motivation

- Growing importance of Large Language Models (LLMs)

- **Challenges in selecting the right LLM:**
  - Increasing number of models
  - Countless evaluation metrics
  - Varying needs

# Dataset Introduction

Hugging Face

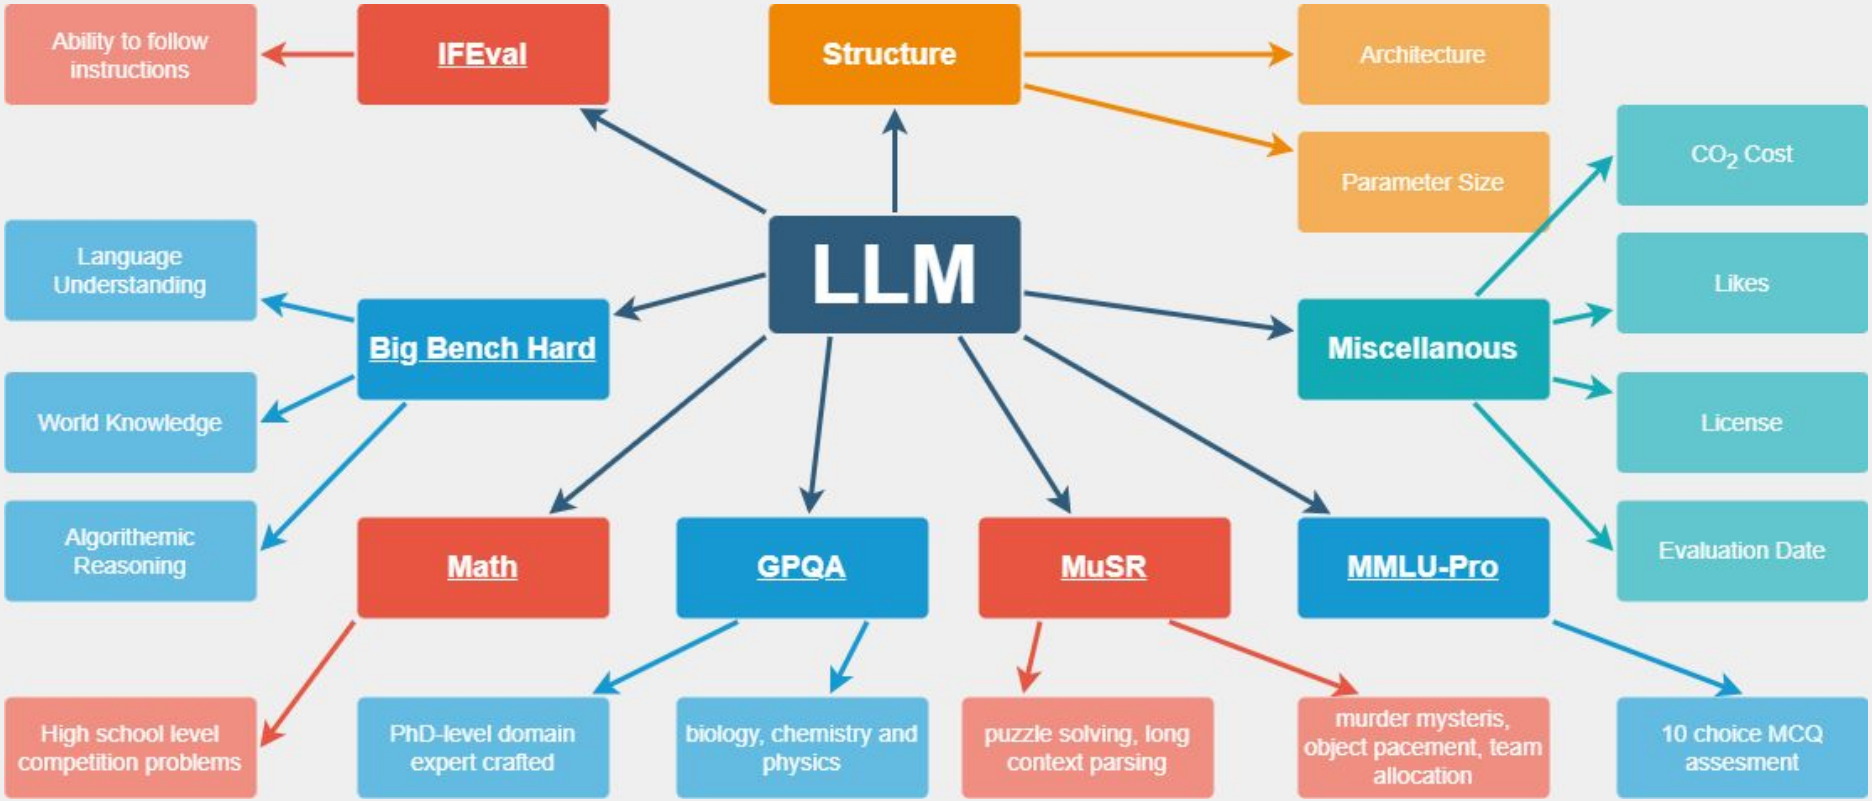

# Methodology


- Correlations
- Importance of Architecture

- Industry requirement
- Ranking models

# Same Performance but Smaller size, Lower Cost

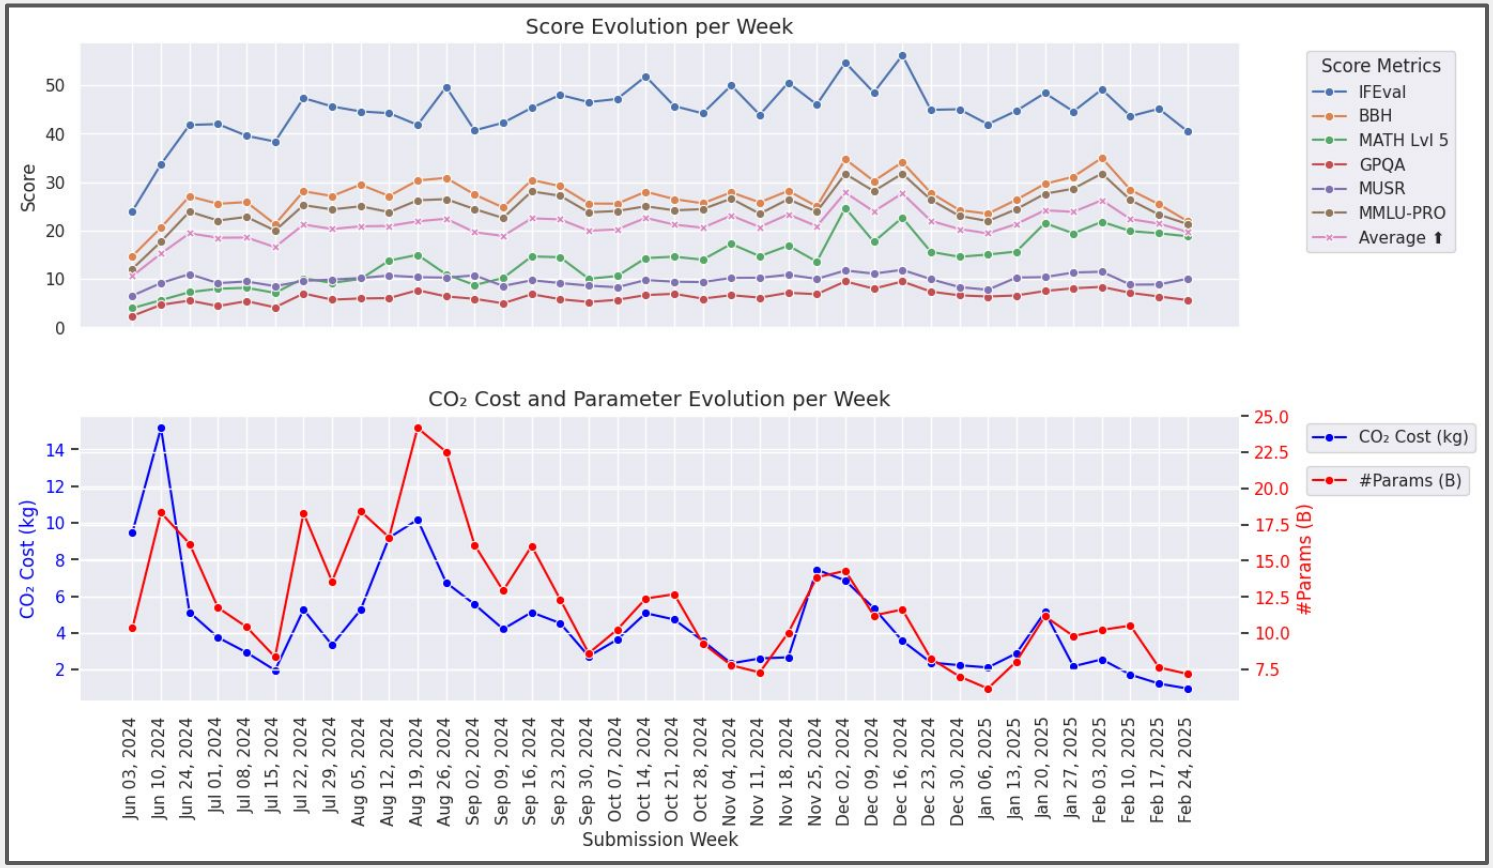

# Efficiency Matters

- Parameter Efficiency
  - Metrics Average / #Params (B)
- Environmental Efficiency
  - Metrics Average / $CO_2$ cost (kg)
- **Architecture** is important

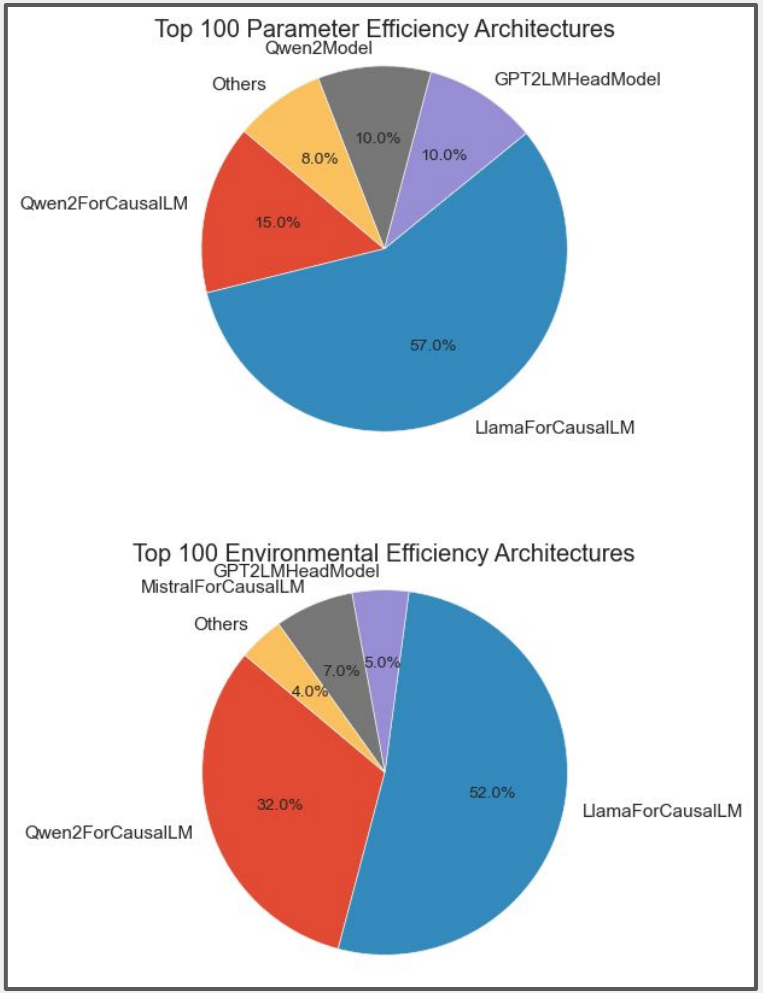

# Architecture- Max Performance

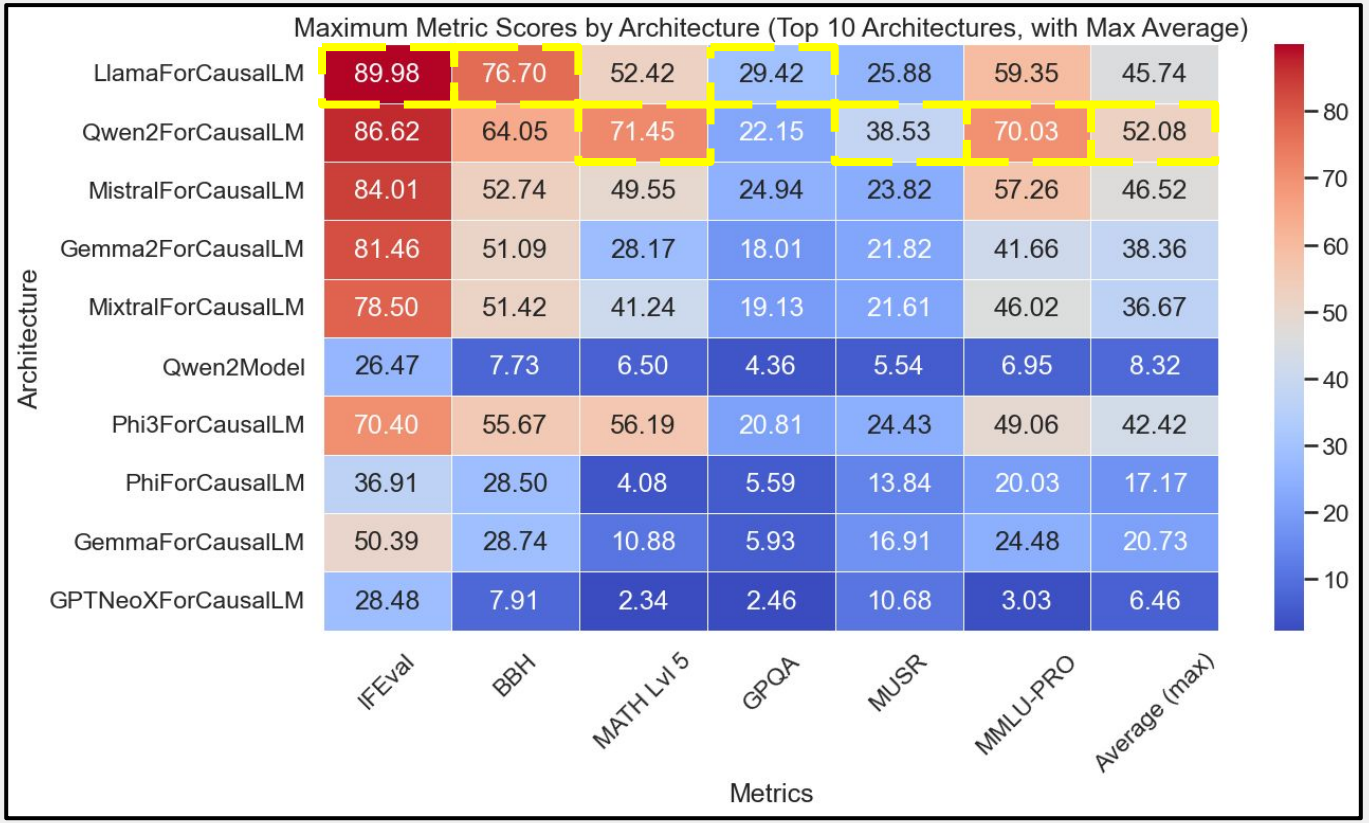

# Architecture - Median Average Score

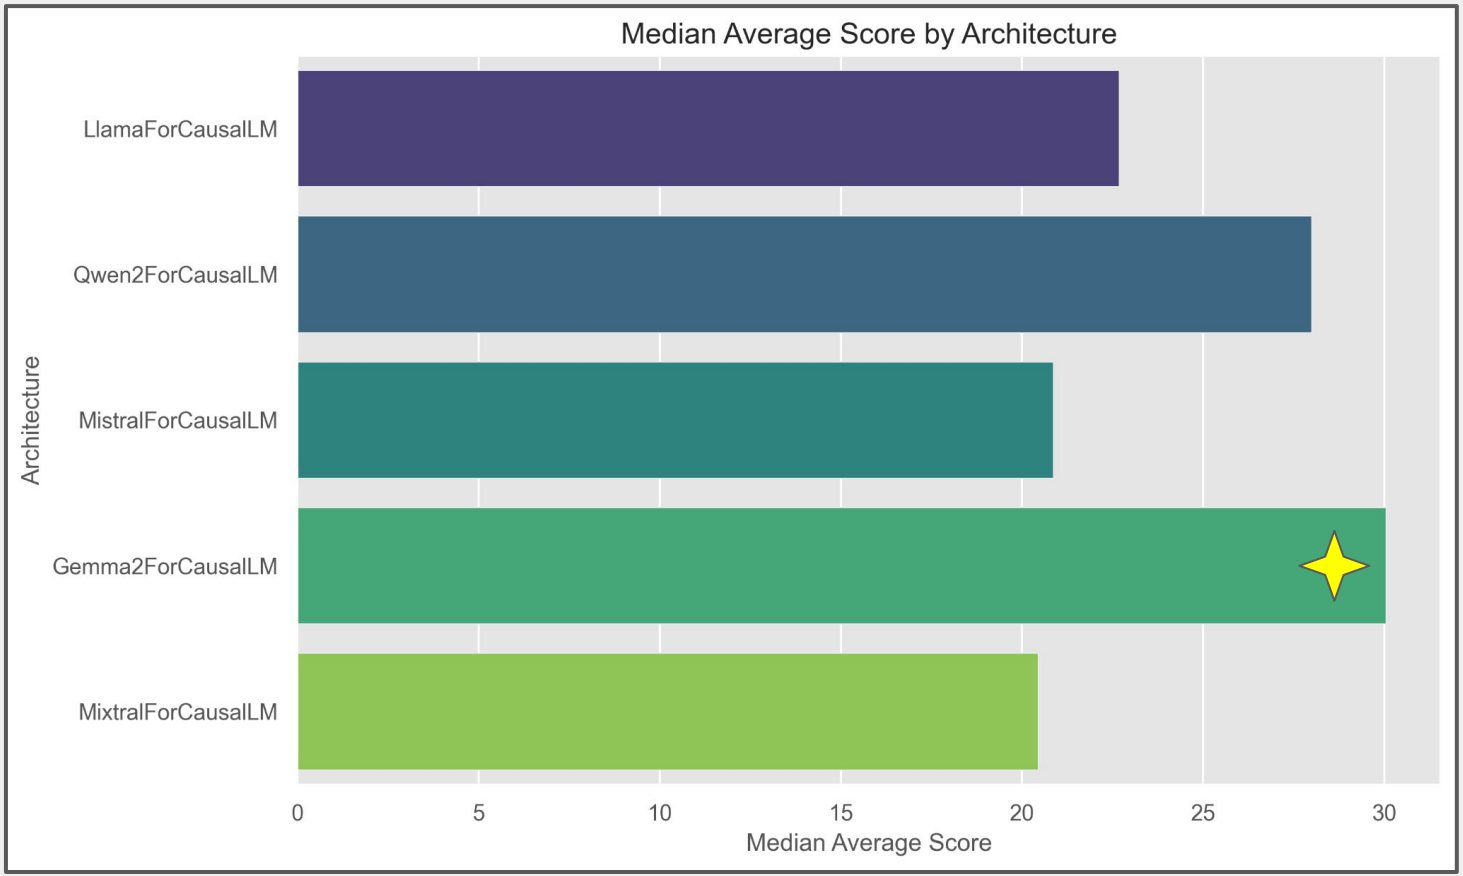

# Architecture - Efficiency Evolution

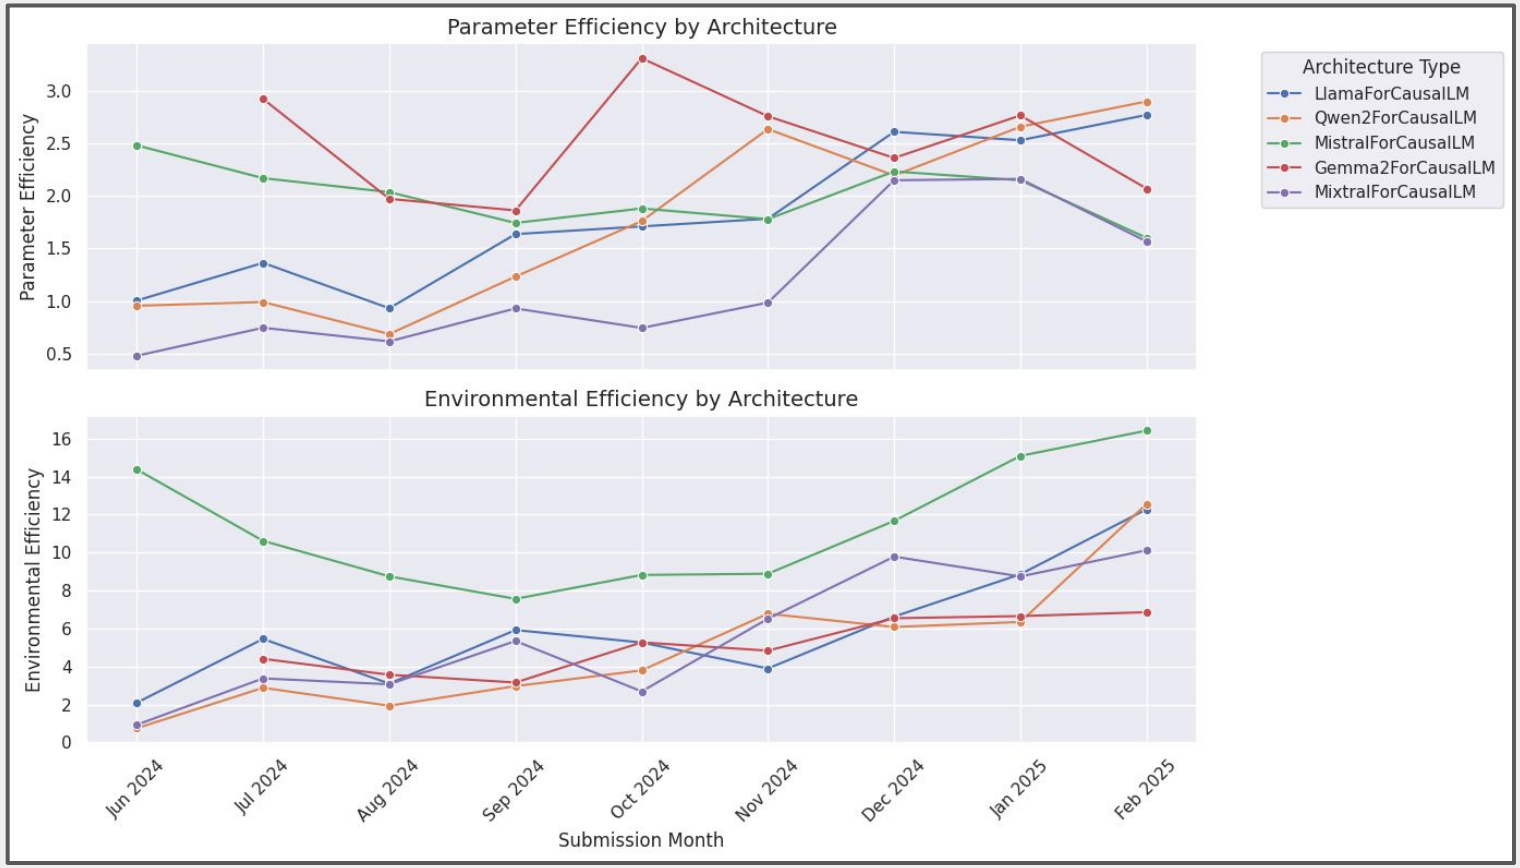

# Industry Requirements

## Tech

- Reasoning
- Solve complex problem

## Academic

- Reasoning
- Ability in maths and science

### Customer Service

- Chatbot
- Relevant Conv.
- Knowledge retention

### Manufacturing

- Engineering proficiency
- Instruction adherence

# Real World Use - Workflow

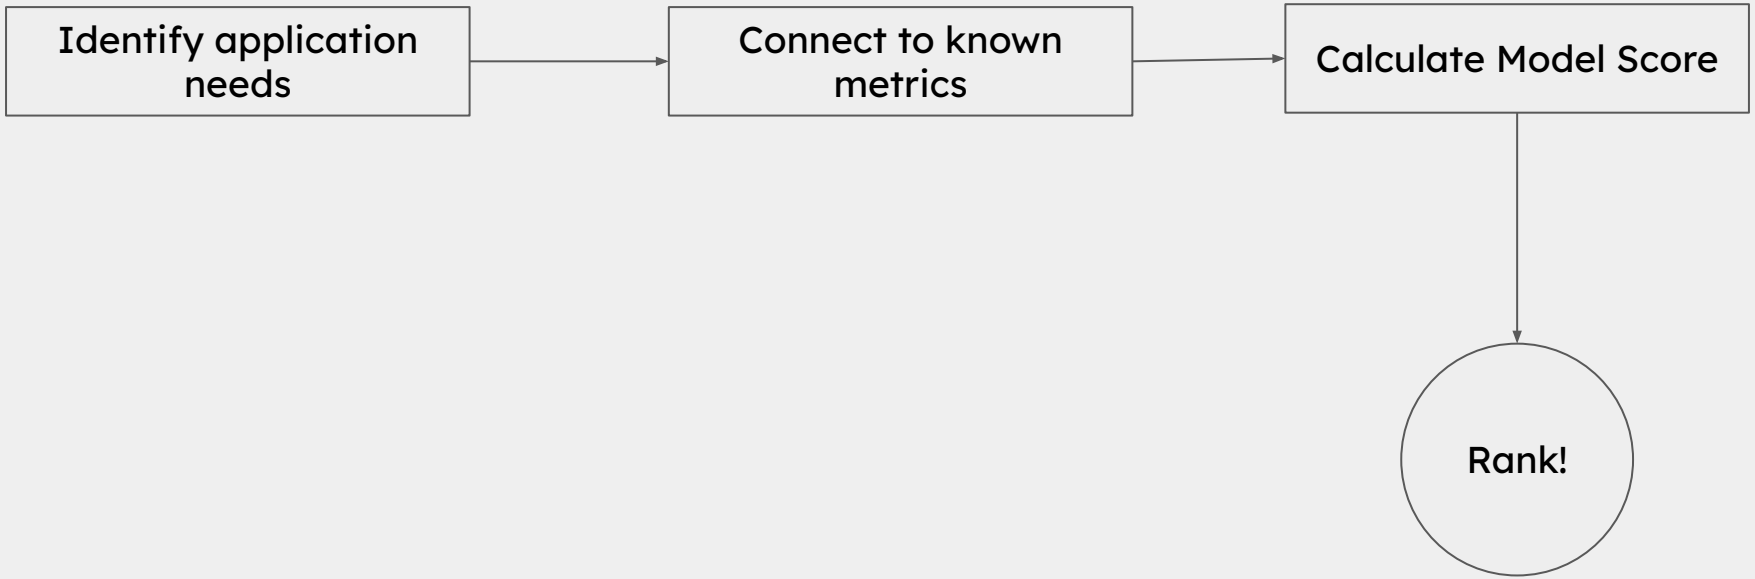

# Real World Use - Linking to Data

**Tech**

- MuSR
- # of Params
- Architecture

**Academic**

- MuSR
- Math lv 5
- GPQA

**Customer Service**

- IFEval
- MMLU-Pro
- Chatbot Template

**Manufacturing**

- MMLU-Pro
- BBH
- IFEval
- Architecture
- Fine-Tune

# Real World Use - Some ideas on ranking

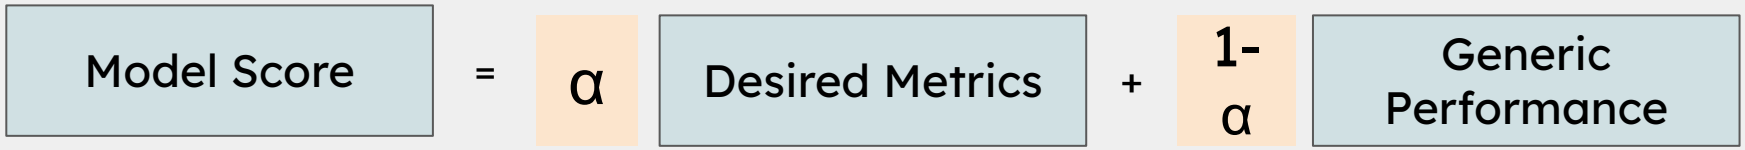


- Desired Metrics
  - Numerical - MMLU-Pro, GPQA
  - Categorical - Architecture, Fine Tuning
- α: A coefficient of confidence (from 0 -1) on Desired Metrics
  - High confidence → Good for specific industry
  - Low confidence → Generally well performing model

# Demo

```
manufac_num_cri = ["MMLU-PRO", "BBH", "IFEval"]
best_architecture_for_manufacturing = [("Architecture", "LlamaForCausalLM"), ("Architecture", "GPTJForCausalLM"), ("Architecture", "CohereForCausalLM"),
                                       ("Architecture", "T5ForConditionalGeneration"), ("Architecture", "RwkvForCausalLM")]
find_tune = [("Type", "🔶 fine-tuned on domain-specific datasets")]
manu_str_cri = best_architecture_for_manufacturing + find_tune

print(get_rank(manufac_num_cri, confident=0.5, num_top_model=3, str_criteria=manu_str_cri, range_matrices = {"CO₂ cost (kg)": (0, 8), "#Params (B)": (0, 10)}))
```

```
                                 eval_name      Architecture
2282     ehristoforu_falcon3-ultraset_float16  LlamaForCausalLM
3875  unsloth_phi-4-unsloth-bnb-4bit_bfloat16  LlamaForCausalLM
3874           unsloth_phi-4-bnb-4bit_bfloat16  LlamaForCausalLM
```

**Input**

- MMLU, BBH, IFEval
- List of architecture
- Fine-tune Model
- < 10 B parameters
- $CO_2$ cost: 0 - 8 kg

**Output**

- Best Models - Falcon

# Demo Results

| Industry Name | Input | Output | Architecture |
|---|---|---|---|
| Tech | "$CO_2$ cost (kg)": (4, 8)<br>"#Params (B)": (20, 30)<br>'MMLU-PRO': (30, 40)<br>["MUSR", "IFEval"] | • Sumatra-20b<br>• Venti-Blend-sce<br>• Venti-20b | • LlamaForCausalLM<br>• LlamaForCausalLM<br>• LlamaForCausalLM |
| Academic | "$CO_2$ cost (kg)": (8, 12)<br>"#Params (B)": (0,35)<br>["MUSR", "MATH Lvl 5", "GPQA"] | • ultiima-32B<br>• Qwen2.5-32B-Instruct<br>• lambda-qwen2.5-32b-dpo-test | • Qwen2ForCausalLM<br>• QwenForCausalLM<br>• QwenForCausalLM |

# Conclusion


Performance Trends
Efficiency Trends
Real World Use
Performance seems to saturate
Architecture's choice determines efficiency
Industry breakdown LLM

# Thanks!

Reference:


- [2407.07000] Etalon: Holistic Performance Evaluation Framework for LLM Inference Systems
- The Gap Between Open and Closed AI Models Might Be Shrinking. Here's Why That Matters
- Large Language Models for Manufacturing
- Why Large Language Models (LLMs) are the future of manufacturing | World Economic Forum
- LLM Chatbot Evaluation Explained: Top Metrics and Testing Techniques - Confident AI